

Laoh

Your portfolio bought AI. *The P&L can't see it.*

A capability diagnostic for the engineering teams inside your portcos - who can actually use AI, what's blocking them, what to fix first.

From the team that screens forward-deployed engineers for frontier AI labs and 3 vertical AI companies.


AI impact has moved from story to diligence requirement.



Most portfolios remain at Deploy - licenses out, operating model unchanged (BCG, 2026).

Vista built an in-house Agentic AI Factory. Mega-funds hire AI operating partners (Korn Ferry). **Mid-market firms rent the playbook.**

Engineering teams sit at six distinct levels of AI capability.



Every level is claimed in interviews. Each one is verifiable only by watching the work - which is the diagnostic.

A team at L₄ ships several times the output of the same team at L₂.

	WHAT GOES IN	WHAT HAPPENS	WHAT COMES OUT
L ₂	one task per person	person and AI write together; every line human-read	one feature at a time; output scales only with headcount
L ₃	three tasks per person	three agents build in parallel; person reviews change summaries	several workstreams move at once; throughput per head multiplies
L ₄	a written spec, tests defined upfront	agents build unsupervised for hours; an independent test suite verifies	shipped software, proven by passing tests; cycle time collapses
L ₅ runs in production today. StrongDM's software factory turns specs into verified software with no human reading code, at ~\$1k/day/engineer in tokens. A lighthouse for where this goes.			

Capacity per head and cycle time are the lines a value-creation plan is priced on.

Developers misjudge their own AI productivity by about 40 points.

+20%

how much faster developers believed AI made them

-19%

how much slower they actually were

Randomized trial, 246 real tasks (METR, 2025)

Every adoption survey in your portfolio carries this same gap. Placing people on the ladder requires watching them work.

We measure capability by observing days of real work.

I

Observability agents go onto laptops and repos

Announced and approved. A few days. Zero engineer hours lost.

2

They record how work happens

Which tasks go to agents, how much agent output gets hand-corrected, whether specs and project briefs exist, how results get verified.

3

We score each person against the level definitions

A level - plus the observed evidence behind it.

THE CALIBRATION

A role-matched live simulation calibrates every score.

The same instrument we use to screen forward-deployed engineers for frontier labs.

Workflow signals only - code never leaves the portco's environment · reports go to the person and the CEO · one 30-min security review with the portco CTO is the full approval

Every person's level comes with the evidence behind it.


A. Mehta

Integrations engineer · Portco A · assessed 12 Jun



operates at L2 - pairs with AI, still reviews every line

SCORED DIMENSIONS



WHAT WE SAW

“hand-edited 40% of the agent's output”

“ran the code once and called it verified”

“knew exactly when not to use the agent - strongest signal”

path to L3: run parallel sessions · stop hand-editing output · build the repo context file

A growth map, never a report card.

Three moves take a portco from baseline to proven capability gains.

I

Measure

2-3 weeks, one portco: capability map, readiness scorecard, sequenced 100-day plan.

2

Move the levels

12-week cohorts grouped by assessed level, practiced on real tickets.

3

Prove it moved


Quarterly re-measure with telemetry on - movement in numbers your IC can read.

ILLUSTRATIVE

hand-corrected agent output ↓ + agent-authored share ↑ → cycle time on comparable tickets ↓ → **capacity per head ↑**

An ugly baseline seen now is pre-diligence on your terms.

We've watched AI deployment fail from the other side of the table.



We screen FDEs for **ElevenLabs · Retell · +3 YC-backed AI cos**

Pulkit Walia

Built the engineering assessment and skilling system at Urban Company (\$3B IPO), a tech-enabled trades marketplace.

Rohit Jain

25,000+ engineers trained. Hundreds of big-tech interview loops run from the other side of the table.

START HERE

Start with one portco - \$1,000 a seat, results in three weeks.

WHAT YOU GET

- 2-3 weeks, one portco
- level map by function, with evidence
- readiness scorecard
- sequenced 100-day plan
- typical run: 30-40 seats, \$30-40k, fund or portco budget

THEN, IF IT EARNS IT

one portco

portfolio dashboard

upskilling / deal diligence

Start where engineering cost is the thesis.

“For \$40k and three weeks we get a measured answer on whether [portco]’s AI spend is building capability or shelfware - before a buyer’s diligence team asks.”

30 mins to scope it. We'll bring the instrument.